

Wizard Pseudo Machine  
Volume One, Programmer's Guide, revision 0.0.10

Tuomo Petteri Venäläinen

April 30, 2013



# Contents

|           |                                   |           |
|-----------|-----------------------------------|-----------|
| <b>I</b>  | <b>Preface</b>                    | <b>5</b>  |
| 1         | Notes                             | 7         |
| 2         | Changes                           | 9         |
| <br>      |                                   |           |
| <b>II</b> | <b>Pseudo Machine</b>             | <b>11</b> |
| <br>      |                                   |           |
| <b>3</b>  | <b>Architecture</b>               | <b>13</b> |
| 3.1       | Memory Map . . . . .              | 14        |
| <br>      |                                   |           |
| <b>4</b>  | <b>Instruction Set</b>            | <b>15</b> |
| 4.1       | Instruction Reference . . . . .   | 15        |
| 4.1.1     | Instructions and Flags . . . . .  | 15        |
| 4.1.2     | Instruction Table . . . . .       | 17        |
| 4.2       | Machine Operations . . . . .      | 18        |
| 4.3       | Reference . . . . .               | 19        |
| 4.3.1     | Opcode Format . . . . .           | 19        |
| 4.3.2     | Instruction Set . . . . .         | 20        |
| 4.3.2.1   | Cheat Sheet . . . . .             | 22        |
| <br>      |                                   |           |
| <b>5</b>  | <b>Assembly</b>                   | <b>23</b> |
| 5.1       | Syntax . . . . .                  | 23        |
| 5.2       | Assembler Directives . . . . .    | 24        |
| 5.2.1     | Input Directives . . . . .        | 24        |
| 5.2.1.1   | .include . . . . .                | 24        |
| 5.2.1.2   | .import . . . . .                 | 25        |
| 5.2.2     | Link Directives . . . . .         | 25        |
| 5.2.2.1   | .org . . . . .                    | 25        |
| 5.2.2.2   | .space . . . . .                  | 25        |
| 5.2.2.3   | .align . . . . .                  | 25        |
| 5.2.2.4   | .globl . . . . .                  | 25        |
| 5.2.3     | Data Directives . . . . .         | 25        |
| 5.2.3.1   | .long . . . . .                   | 25        |
| 5.2.3.2   | .byte . . . . .                   | 25        |
| 5.2.3.3   | .short . . . . .                  | 26        |
| 5.2.3.4   | .asciz . . . . .                  | 26        |
| 5.2.4     | Preprocessor Directives . . . . . | 26        |

|         |                            |    |
|---------|----------------------------|----|
| 5.2.4.1 | .define                    | 26 |
| 5.3     | Input and Output           | 27 |
| 5.3.1   | Simple Program             | 27 |
| 5.4     | Threads                    | 27 |
| 5.4.1   | Example Program            | 28 |
| 5.5     | Hooks                      | 28 |
| 5.5.1   | Pre-Defined Hooks          | 28 |
| 5.5.2   | Hook Interface             | 28 |
| 5.5.3   | Example Program            | 28 |
| 5.6     | Interrupts                 | 29 |
| 5.6.1   | Break Points               | 29 |
| 5.6.2   | Interrupt Interface        | 30 |
| 5.6.3   | Keyboard Input             | 30 |
| 5.6.3.1 | Keyboard Interrupt Handler | 30 |
| 5.6.3.2 | Keyboard Support Code      | 30 |

## Part I

# Preface



# Chapter 1

## Notes

### **Brief**

Wizard Pseudo Machine, WPM, is a software-based virtual machine. It can be programmed in its own instruction set / assembly dialect, using AT&T syntax assembly. The incorporated assembler (**zas**) does not yet support loading pre-translated objects, but translates assembly code into bytecode before execution. This bytecode is then handled to the pseudo machine (**wpm**) to be executed.

The pseudo machine instruction set is designed to be a close match to basic operations of the C programming language, with some other typical machine operations.

### **Background and Future**

Wizard Pseudo Machine project started as an attempt to create a tool for educational purposes. The machine is a virtual processor with custom instruction set that is close to existing ones; goals include making this instruction set simple, relatively complete, and useful for C language operations in case we decide to create a C compiler or code generator for an existing one.

At the time of this writing, the virtual assembler is in good shape. The assembler provides not only `.include` to place other files into the stream verbatim, but also `.import` to 'link' with other assembly files with access to their global symbols.

The project is by no means complete yet. I wish for useful libraries for audio and graphics, a useful stock 'standard' library, and lots of other things to happen in the near future. :)





# Chapter 2

## Changes

### 0.0.10

- added section **Instruction Reference** along with CPU-flag management by some instructions
- added PDF-index and hyperlink-TOC

### 0.0.9

- grammatical and layout corrections

### 0.0.8

- added .asciz
- added new instruction **hook** to invoke high level services; this lets us run them with native code in the virtual machine

### 0.0.7

- adding more content
- cleaning up

### 0.0.6

- still fixing typos
- added a new **Architecture** Chapter

### 0.0.5

- Added text on threading as well as a couple of assembly examples

### 0.0.4

- changed the title, made a couple of mistakes there; the book is now correctly called **Wizard Pseudo Machine**

### 0.0.3

- added new subsection Opcode Format

**0.0.2**

- reorganised some assembly sections; added `.space`, `.long`, `.short`, and `.byte`

**0.0.1**

- changed the assembler to use `.include` and `.import` instead of `#include` and `#import`
- changed the term 'argument' to 'operand' in many places in this booklet

## Part II

# Pseudo Machine



## Chapter 3

# Architecture

Wizard Pseudo Machine is an architecture with 32-bit machine words; however, room has been left in the implementation for 64-bit support.

The pseudo machine supports flat 4-gigabyte address space, of which some is mapped for interrupt vector, interrupt handlers, thread and interrupt stacks, graphics, and other purposes.

Native word size is **32-bit**. Words are little endian, i.e. lowest byte first.

There exists an instruction, **thr**, to start executing new threads from desired code locations in memory.

We follow the **von Neumann architecture**, so we basically have 3 abstractions; **CPU**, **memory**, and **input/output**.

The machine is a purposefully **RISC-like load-store** approach, meaning there is only a single load-store instruction (**mov**) that deals with memory-addressed operands.

### 3.1 Memory Map

#### Notes

| Address         | Purpose          | Brief   |
|-----------------|------------------|---|
| 0               | interrupt vector | interrupt handler descriptors                     |
| 4096            | keyboard buffer  | keyboard input queue                              |
| 8192            | text segment     | application program code (read-execute)           |
| 8192 + TEXTSIZE | data segment     | program data (read-write)                         |
| DATA + DATASIZE | BSS segment      | uninitialised data (runtime-allocated and zeroed) |
| MEMSIZE - 3.5 G | dynamic segment  | free space for slab allocator                     |
| 3.5 gigabytes   | graphics         | 32-bit ARGB-format draw buffer                    |

- the VM's memory size is currently specified as **MEMSIZE**
- thread stacks live at **MEMSIZE - thr\_id \* THRSTKSIZE**, i.e. at top of 'physical' address space

## Chapter 4

# Instruction Set

The WPM instruction set was designed to resemble the C language closely, as well as support a RISC-oriented set of typical machine operations.

### 4.1 Instruction Reference

#### 4.1.1 Instructions and Flags

##### Operands

- **i** stands for immediate operand
- **r** stands for register operand
- **m** stands for memory operand

##### Flags

Certain instructions set bits in the machine status word (MSW). This is documented here on per-instruction basis.

- **z** stands for zero flag (ZF)
- **c** stands for carry flag (CF)
- **o** stands for overflow flag (OF)
- **s** stands for sign flag (SF)

**TODO:** stack/call conventions for certain instructions





## 4.1.1.1 Instruction Table

| Mnemonic       | Operand #1 | Operand #2 | Brief                                   | Flags   |
|----------------|------------|------------|---|---------|
| <b>not</b>     | r          | N/A        | op1 = op1;                              | z       |
| <b>and</b>     | r, i       | r          | op2 = op2 & op1;                        | z       |
| <b>or</b>      | r, i       | r          | op2 = op2   op1;                        | N/A     |
| <b>xor</b>     | r, i       | r          | op2 = op2 ^ op1;                        | z       |
| <b>shr</b>     | r, i       | r          | op2 = op2 » op1; (zero-fill)            | z       |
| <b>shra</b>    | r, i       | r          | op2 = op2 » op1; (sign-fill)            | z       |
| <b>shl</b>     | r, i       | r          | op2 = op2 « op1;                        | o, c    |
| <b>ror</b>     | r, i       | r          | op2 = op2 ROR op1;                      | c       |
| <b>rol</b>     | r, i       | r          | op2 = op2 ROL op1;                      | c       |
| <b>inc</b>     | r, i       | N/A        | op1++;                                  | o       |
| <b>dec</b>     | r, i       | N/A        | op1--;                                  | o, z    |
| <b>add</b>     | r, i       | r          | op2 = op2 + op1;                        | o, z    |
| <b>sub</b>     | r, i       | r          | op2 = op2 - op1;                        | s, z    |
| <b>cmp</b>     | r, i       | r          | compare two values and set flags        | s, z    |
| <b>mul</b>     | r, i       | r          | op2 = op2 * op1;                        | o, s, z |
| <b>div</b>     | r, i       | r          | op2 = op2 / op1;                        | s, z    |
| <b>mod</b>     | r, i       | r          | op2 = op2 % op1;                        | N/A     |
| <b>bz</b>      | r, i       | N/A        | branch to op1 if (CF == 0)              | N/A     |
| <b>bnz</b>     | r, i       | N/A        | branch to op1 if (CF != 0)              | N/A     |
| <b>blt</b>     | r, i       | N/A        | branch to op1 if (SF != 0)              | N/A     |
| <b>ble</b>     | r, i       | N/A        | branch to op1 if (SF != 0)    (ZF == 0) | N/A     |
| <b>bgt</b>     | r, i       | N/A        | branch to op1 if (SF != 0) && (ZF != 0) | N/A     |
| <b>bge</b>     | r, i       | N/A        | branch to op1 if (SF != 0)    (ZF == 0) | N/A     |
| <b>bo</b>      | r, i       | N/A        | branch to op1 if (OF != 0)              | N/A     |
| <b>bno</b>     | r, i       | N/A        | branch to op1 if (OF == 0)              | N/A     |
| <b>bc</b>      | r, i       | N/A        | branch to op1 if (CF != 0)              | N/A     |
| <b>bnc</b>     | r, i       | N/A        | branch to op1 if (CF == 0)              | N/A     |
| <b>pop</b>     | N/A        | N/A        | pop top of stack                        | N/A     |
| <b>push</b>    | r, i       | N/A        | push value on stack                     | N/A     |
| <b>mov</b>     | r, i, m    | r, i, m    | load or store 32-bit longword           | N/A     |
| <b>movb</b>    | r, i, m    | r, i, m    | load or store 8-bit byte                | N/A     |
| <b>movw</b>    | r, i, m    | r, i, m    | load or store 16-bit word               | N/A     |
| <b>call</b>    | r, i       | N/A        | call subroutine; construct stack frame  | N/A     |
| <b>enter</b>   | N/A        | N/A        | enter subroutine                        | N/A     |
| <b>leave</b>   | N/A        | N/A        | leave subroutine                        | N/A     |
| <b>ret</b>     | N/A        | N/A        | return from subroutine                  | N/A     |
| <b>lmsw</b>    | r, i       | N/A        | load machine status word                | N/A     |
| <b>smsw</b>    | m          | N/A        | load machine status word                | N/A     |
| <b>reset</b>   | m          | N/A        | reset machine                           | N/A     |
| <b>hlt</b>     | m          | N/A        | halt machine                            | N/A     |
| <b>nop</b>     | N/A        | N/A        | no/dummy operation                      | N/A     |
| <b>brk</b>     | N/A        | N/A        | breakpoint                              | N/A     |
| <b>trap</b>    | N/A        | N/A        | software interrupt                      | N/A     |
| <b>cli</b>     | N/A        | N/A        | disable interrupts                      | N/A     |
| <b>sti</b>     | N/A        | N/A        | enable interrupts                       | N/A     |
| <b>iret</b>    | N/A        | N/A        | return from interrupt handler           | N/A     |
| <b>thr</b>     | m          | N/A        | start new thread at address             | N/A     |
| <b>cmpswap</b> | m          | r, i       | atomic compare and swap                 | N/A     |
| <b>inb</b>     | r, i       | N/A        | read 8-bit byte from input port         | N/A     |
| <b>outb</b>    | r, i       | N/A        | write 8-bit byte to input port          | N/A     |
| <b>inw</b>     | r, i       | N/A        | read 16-bit word from input port        | N/A     |
| <b>outw</b>    | r, i       | N/A        | write 16-bit word to input port         | N/A     |
| <b>inl</b>     | r, i       | N/A        | read 32-bit longword from input port    | N/A     |
| <b>outl</b>    | r, i       | N/A        | write 32-bit longword to input port     | N/A     |
| <b>hook</b>    | r, i       | N/A        | invoke system service                   | N/A     |

## 4.2 Machine Operations

The following is a C-code snippet listing machine instructions and their IDs in opcodes.

```
/* instruction set */

/* ALU instructions */
#define OPNOT      0x01 // 2's complement, args
#define OPAND      0x02 // logical AND
#define OPOR       0x03 // logical OR
#define OPXOR      0x04 // logical exclusive OR
#define OPSHR      0x05 // logical shift right (fill with zero)
#define OPSHRA     0x06 // arithmetic shift right (fill with sign)
#define OPSHL      0x07 // shift left (fill with zero)
#define OPROR      0x08 // rotate right
#define OPROL      0x09 // rotate left
#define OPINC      0x0a // increment by one
#define OPDEC      0x0b // decrement by one
#define OPADD      0x0c // addition
#define OPSUB      0x0d // subtraction
#define OPCMP      0x0e // compare
#define OPMUL      0x0f // multiplication
#define OPDIV      0x10 // division
#define OPMOD      0x11 // modulus
#define OPBZ       0x12 // branch if zero
#define OPBNZ      0x13 // branch if not zero
#define OPBLT      0x14 // branch if less than
#define OPBLE      0x15 // branch if less than or equal to
#define OPBGT      0x16 // branch if greater than
#define OPBGE      0x17 // branch if greater than or equal to
#define OPBO       0x18 // branch if overflow
#define OPBNO      0x19 // branch if no overflow
#define OPBC       0x1a // branch if carry
#define OPBNC      0x1b // branch if no carry
#define OPPOP      0x1c // pop from stack
#define OPPUSH     0x1d // push to stack
#define OPMOV      0x1e // load/store 32-bit longword
#define OPMOVB     0x1f // load/store 8-bit byte
#define OPMOVW     0x20 // load/store 16-bit word
#define OPJMP      0x21 // jump to given address
#define OPCALL     0x22 // call subroutine
#define OPENTER    0x23 // subroutine prologue
#define OPLEAVE    0x24 // subroutine epilogue
#define OPRET      0x25 // return from subroutine
#define OPLMSW     0x26 // load machine status word
#define OPSMSW     0x27 // store machine status word
#define OPRESET    0x28 // reset into well-known state
#define OPNOP      0x29 // dummy operation
#define OPHLT      0x2a // halt execution
```

```

#define OPBRK      0x2b // breakpoint
#define OPTRAP     0x2c // trigger a trap (software interrupt)
#define OPCLI      0x2d // disable interrupts
#define OPSTI      0x2e // enable interrupts
#define OPIRET     0x2f // return from interrupt handler
#define OPTHR      0x30 // start new thread at given address
#define OPCMP_SWAP 0x31 // atomic compare and swap
#define OPINB      0x32 // read 8-bit byte from port
#define OPOUTB     0x33 // write 8-bit byte to port
#define OPINW      0x34 // read 16-bit word
#define OPOUTW     0x35 // write 16-bit word
#define OPINL      0x36 // read 32-bit long
#define OPOUTL     0x37 // write 32-bit long
#define OPHOOK     0x38 // system services

```

## 4.3 Reference

### 4.3.1 Opcode Format

The following C structure is what the stock assembler uses for opcode output.

#### Opcode Structure

```

struct wpmopcode {
    unsigned inst      : 8; // instruction ID
    unsigned unit      : 2; // unit ID
    unsigned arg1t     : 3; // argument #1 type
    unsigned arg2t     : 3; // argument #2 type
    unsigned reg1      : 6; // register #1 ID + addressing flags
    unsigned reg2      : 6; // register #2 ID + addressing flags
    unsigned size      : 2; // 1..3, shift count
    unsigned res       : 2; // reserved
    int32_t args[2];
} __attribute__((packed));

```

#### Notes

- **inst** is the instruction ID; 0 is invalid
- **unit** is a future unit ID; ALU, FPU, VPU (SIMD), GPU?
- **reg1** and **reg2** are source and destination register IDs
- **operation size** can be calculated as **op->size** « **2**
- **res**-bits are reserved for future extensions
- **args** contains 0, 1, or 2 32-byte addresses or values

### 4.3.2 Instruction Set

#### Operand Types

- **r** stands for register operand
- **i** stands for immediate operand value
- **a** stands for immediate direct address operand
- **p** stands for indirect address operand
- **n** stands for indexed address operand
- **m** stands for all of **a**, **i**, and **n**

#### Notes

- C language doesn't specify whether right shifts are arithmetic or logical
- Arithmetic right shift fills leftmost 'new' bits with the sign bit, logical shift fills with zero; left shifts are always fill rightmost bits with zero

#### Instructions

Below, I will list machine instructions and illustrate their relation to C.

#### Notes

- the **inb()** and other functions dealing with I/O are usually declared through `<sys/io.h>`



## 4.3.2.1 Cheat Sheet

| C Operation  | Instruction | Operands            | Brief                           |
|--------------|-------------|---------------------|---------------------------------|
|              | not         | r dest              | reverse all bits                |
| &            | and         | r/i src, r dest     | logical AND                     |
|              | or          | r/i src, r dest     | logical OR                      |
| ^            | xor         | r/i src, r dest     | logical exclusive OR            |
| «            | shl         | r/i cnt, r dest     | shift left by count             |
| »            | shr         | r/i cnt, r dest     | arithmetic shift right          |
|              | shrl        | r/i cnt, r dest     | logical shift right             |
| N/A          | ror         | r/i cnt, r dest     | rotate right by count           |
| N/A          | rol         | r/i cnt, r dest     | rotate left by count            |
| ++           | inc         | r dest              | increment by one                |
| -            | dec         | r dest              | decrement by one                |
| +            | add         | r/i cnt, r dest     | addition                        |
| -            | sub         | r/i cnt, r dest     | subtraction                     |
| ==, != etc.  | cmp         | r/i src, r dest     | comparison; sets MSW-flags      |
| *            | mul         | r/i src, r dest     | multiplication                  |
| /            | div         | r/i src, r dest     | division                        |
| %            | mod         | r/i src, r dest     | modulus                         |
| ==, !        | bz          | none                | branch if zero                  |
| !=, (val)    | bnz         | none                | branch if not zero              |
| <            | blt         | none                | branch if less than             |
| <=           | ble         | none                | branch if less than or equal    |
| >            | bgt         | none                | branch if greater than          |
| >=           | bge         | none                | branch if greater than or equal |
| N/A          | bo          | none                | branch if overflow              |
| N/A          | bno         | none                | branch if no overflow           |
| N/A          | bc          | none                | branch if carry                 |
| N/A          | bnc         | none                | branch if no carry              |
| dest = *sp++ | pop         | r dest              | pop from stack                  |
| -sp = src    | push        | r src               | push onto stack                 |
| dest = src   | mov         | r/i/m src, r/m dest | load/store longword (32-bit)    |
| dest = src   | movb        | r/i/m src, r/m dest | load/store byte (8-bit)         |
| dest = src   | movw        | r/i/m src, r/m dest | load/store word (16-bit)        |
| N/A          | jmp         | r/m dest            | continue execution at dest      |
| N/A          | call        | a/p dest            | call subroutine                 |
| N/A          | enter       | none                | subroutine prologue             |
| N/A          | leave       | none                | subroutine epilogue             |
| N/A          | ret         | none                | return from subroutine          |
| N/A          | lmsw        | r/i dest            | load machine status word        |
| N/A          | smsw        | r/i src             | store machine status word       |
| N/A          | reset       | none                | reset machine                   |
| N/A          | nop         | none                | no operation                    |
| N/A          | hlt         | none                | halt machine                    |
| N/A          | brk         | none                | breakpoint                      |
| N/A          | trap        | r/i src             | trigger software interrupt      |
| N/A          | cli         | none                | disable interrupts              |
| N/A          | sti         | none                | enable interrupts               |
| N/A          | iret        | none                | return from interrupt handler   |
| N/A          | thr         | r/i dest            | start thread at dest            |
| N/A          | cmpswap     | r/i src, m dest     | atomic compare and swap         |
| inb()        | inb         | r/i src             | read byte (8-bit)               |
| outb()       | outb        | r/i src, r/i dest   | write byte (8-bit)              |
| inw()        | inw         | r/i src             | read word (16-bit)              |
| outw()       | outw        | r/i src, r/i dest   | write word (16-bit)             |
| inl()        | inl         | r/i src             | read longword (32-bit)          |
| outl()       | outl        | r/i src, r/i dest   | write longword (32-bit)         |
| hook()       | hook        | r1 cmd, r2 arg      | system service                  |

# Chapter 5

## Assembly

### 5.1 Syntax

#### AT&T Syntax

We use so-called AT&T-syntax assembly. Perhaps the most notorious difference from Intel-syntax is the operand order; AT&T lists the source operand first, destination second, whereas Intel syntax does it vice versa.

#### Symbol Names

Label names must start with an underscore or a letter; after that, the name may contain underscores, letters, and digits. Label names end with a ':', so like

```
value: .long 0xb4b5b6b7
```

would declare a longword value at the address of **value**.

#### Instructions

The instruction operand order is source first, then destination. For example,

```
mov 8(%r0), val
```

would store the value from address **r0** + **8** to the address of the label **val**.

#### Operands

Register operand names are prefixed with a '%'. Immediate constants and direct addresses are prefixed with a '#'. Label addresses are referred to as their names without prefixes.

The assembler supports simple preprocessing (of constant-value expressions), so it is possible to do things such as

```
.define FLAG1 0x01  
.define FLAG2 0x02
```

```
mov $(FLAG1| FLAG2), %r1
```

### Registers

Register names are prefixed with '%'; there are 16 registers r0..r15. For example,

```
add %r0, %r1
```

would add the longword in r0 to r1.

### Direct Addressing

Direct addressing takes the syntax

```
mov val, %r0
```

which moves the longword at **address val** into r0.

### Indexed Addressing

Indexed addressing takes the syntax

```
mov 4(%r0), %r1
```

where 4 is an integral constant offset and r0 is a register name. In short, this would store the value at the address **r0 + 4** into r1.

### Indirect Addressing

Indirect addresses are indicated with a '\*', so

```
mov *%r0, %r1
```

would store the value from the **address in the register r0** into register r1, whereas

```
mov *val, %r0
```

would move the value **pointed to by val** into r0.

Note that the first example above was functionally equivalent with

```
mov (%r0), %r1
```

### Immediate Addressing

Immediate addressing takes the syntax

```
mov $str, %r0
```

which would store the **address of str** into r0.

## 5.2 Assembler Directives

### 5.2.1 Input Directives

#### 5.2.1.1 .include

The **.include** directive takes the syntax

```
.include <file.asm>
```



to insert file.asm into the translation stream verbatim.

#### 5.2.1.2 `.import`

The `.import` directive takes the syntax

```
.import <file.asm>
```

or

```
.import <file.obj>
```

to import foreign assembly or object files into the stream. **Note** that only symbols declared with `.globl` will be made globally visible to avoid namespace pollution.

### 5.2.2 Link Directives

#### 5.2.2.1 `.org`

The `.org` directive takes a single argument and sets the linker location address to the given value.

#### 5.2.2.2 `.space`

The `.space` directive takes a single argument and advances the link location address by the given value.

#### 5.2.2.3 `.align`

The `.align` directive takes a single argument and aligns the next label, data, or instruction to a boundary of the given size.

#### 5.2.2.4 `.globl`

The `.globl` directive takes one or several symbol names arguments and declares the symbols to have global visibility (linkage).

### 5.2.3 Data Directives

#### 5.2.3.1 `.long`

`.long` takes any number of arguments and declares in-memory 32-bit entities.

#### 5.2.3.2 `.byte`

`.byte` takes any number of arguments and declares in-memory 8-bit entities.

### 5.2.3.3 `.short`

`.short` takes any number of arguments and declares in-memory 16-bit entities.

### 5.2.3.4 `.asciz`

`.asciz` takes a C-style string argument of characters enclosed within double quotes (`"`). Escape sequences `'\n'` (newline), `'\t'` (tabulator), and `'\r'` (carriage return) are supported.

## 5.2.4 Preprocessor Directives

### 5.2.4.1 `.define`

`.define` lets one declare symbolic names for constant (numeric) values. For example, if you have

`<hook.def>`

```
.define PZERO 0
.define PALLOC 1
.define PFREE 2
```

you can then use the symbolic names like

```
.include <hook.def>
.import <bzero.asm>

memalloc:
    mov    $16384, %r0
    hook   $PALLOC
    mov    %r0, ptr
    ret

memzero:
    mov    ptr, %r0
    mov    $4096, %r1
    hook   $PZERO
    ret

memfree:
    mov    ptr, %r0
    hook   $PFREE
    ret

_start:
    call   memalloc
    call   memzero
    call   memfree
    hlt
```

```
ptr:    .long    0x00000000
```

## 5.3 Input and Output

The pseudo machine uses some predefined ports for keyboard and console I/O. The currently predefined ports are

| Port | Use            | Notes                          |
|------|----------------|--------------------------------|
| 0x00 | keyboard input | interrupt-driven               |
| 0x01 | console output | byte stream                    |
| 0x02 | error output   | directed to console by default |

### 5.3.1 Simple Program

The following code snippet prints the string "hello" + a newline to the console. Note that the string is saved using the standard C convention of NUL-character termination.

```
msg:    .asciz "hello\n"

.align  4

_start:
    mov    $msg, %r0
    movb   *%r0, %r1
    mov    $0x01, %r2
    cmp    $0x00, %r1
    bz     done
loop:
    inc    %r0
    outb   %r1, %r2
    movb   *%r0, %r1
    cmp    $0x00, %r1
    bnz    loop
done:
    hlt
```

## 5.4 Threads

Wizard Pseudo Machine supports hardware threads with the **thr** instruction. It takes a single argument, which specifies the new execution start address; function arguments should be passed in registers.

### 5.4.1 Example Program

The following piece of code shows simple utilisation of threads.

```
.import <bzero.asm>

memzero:
    mov     $65536, %r0        // address
    mov     $4096, %r1         // length
    call    bzero
    hlt

_start:
    thr     $memzero
    hlt
```

## 5.5 Hooks

Hooks exist to provide system services. Hooks invoke native code in the virtual machine to do things such as manage memory and I/O.

### 5.5.1 Pre-Defined Hooks

| Number | Name   | Purpose   |
|--------|--------|---|
| 0x00   | PZERO  | zero given number of pages at given address         |
| 0x01   | PALLOC | allocate given number of bytes from dynamic segment |
| 0x01   | PFREE  | free region at given address                        |

### 5.5.2 Hook Interface

Hook **parameters** are passed **in registers**. Hook **return value** is stored **in r0**. Here is the current interface definition.

- PZERO takes two arguments; destination address in r0, and region size (in bytes) in r1. PZERO returns nothing.
- PALLOC takes one argument; allocation size in r0. PALLOC returns allocated address or zero on failure.
- PFREE takes one argument; allocation address in r0. PFREE returns nothing.

### 5.5.3 Example Program

The following programs uses hooks to accomplish 3 tasks: allocate 16384 bytes of memory, zero it, and finally free it. In reality this alone is useless, but it serves as an example.

```

.import <bzero.asm>

alloc:
mov     $16384, %r0
hook    $1
mov     %r0, ptr
ret

zero:
mov     ptr, %r0
mov     $16384, %r1
hook    $0
ret

free:
mov     ptr, %r0
hook    $2
ret

_start:
call    alloc
call    zero
call    free
hlt

ptr:     .long    0x00000000

_foo:    .space   4096, 0xff

```

## 5.6 Interrupts

Software- and CPU-generated interrupts are often referred to as **traps**. I call those and hardware-generated **interrupt requests** interrupts, collectively.

### 5.6.1 Break Points

The **brk** instruction triggers a breakpoint interrupt. The default action is to print a stack trace and continue execution.

The **use** of **brk** is simple; just use the zero-operand instruction in your assembly file:

```
brk ; trigger breakpoint
```

### 5.6.2 Interrupt Interface

The lowest page (4096 bytes) in virtual machine address space contains the **interrupt vector**, i.e. a table of interrupt handler addresses to trigger them.

Interrupt handler invocations only push the **program counter** and **old frame pointer**, so you need to reserve the registers you use manually. This is so interrupts could be as little overhead as possible to handle.

### 5.6.3 Keyboard Input

In order to read keyboard input without polling, we need to hook the **interrupt 0**. This is done in two code modules; an interrupt handler as well as other support code.

I will illustrate the interrupt handler first.

#### 5.6.3.1 Keyboard Interrupt Handler

**TODO:** example interrupt handler

#### 5.6.3.2 Keyboard Support Code

**TODO:** queue keypresses in 16-bit values; 32-bit if full Unicode requested.