

# Zero Pseudo Machine

Volume One, Programmer's Guide, revision 0.0.1

Tuomo Petteri Venäläinen

June 24, 2017



# Contents

<b>I</b>	<b>Preface</b>	<b>5</b>
<b>1</b>	<b>Notes</b>	<b>7</b>
<b>II</b>	<b>Pseudo Machine</b>	<b>9</b>
<b>2</b>	<b>Architecture</b>	<b>11</b>
2.1	Memory Organization . . . . .	12
<b>3</b>	<b>Instruction Set</b>	<b>13</b>
3.1	Instruction Reference . . . . .	13
3.1.1	Instruction Set . . . . .	13
3.1.2	Instruction Table . . . . .	15
3.1.3	I/O Port Map . . . . .	16
<b>4</b>	<b>Assembly</b>	<b>17</b>
4.1	Syntax . . . . .	17
4.2	Assembler Directives . . . . .	18
4.3	Input Directives . . . . .	18
4.3.1	.include . . . . .	18
4.3.2	.import . . . . .	19
4.4	Link Directives . . . . .	19
4.4.1	.org . . . . .	19
4.4.2	.space . . . . .	19
4.4.3	.align . . . . .	19
4.4.4	.globl . . . . .	19
4.5	Data Directives . . . . .	19
4.5.1	.long . . . . .	19
4.5.2	.byte . . . . .	20
4.5.3	.short . . . . .	20
4.5.4	.asciz . . . . .	20
4.5.5	Preprocessor Directives . . . . .	20
4.6	.define . . . . .	20
4.7	Input and Output . . . . .	20
4.8	Simple Program . . . . .	20
4.9	Threads . . . . .	21
4.10	Example Program . . . . .	21
4.11	Interrupts . . . . .	21

4.11.1	Interrupt Interface . . . . .	22
4.11.2	Keyboard Input . . . . .	22
4.11.3	Keyboard Interrupt Handler . . . . .	22
4.11.3.1	Keyboard Support Code . . . . .	22

## Part I

# Preface



# Chapter 1

## Notes

### **Brief**

Zero Pseudo Machine, ZPM, is a software-based virtual machine. The machine is programmed in its own assembly dialect; the instruction set is reminiscent of many current RISC-like implementations.





## Part II

# Pseudo Machine



## Chapter 2

# Architecture

### Notes

ZPM is an architecture with 32-bit machine words; room has been left in the implementation for 64-bit support.

ZPM words are little-endian (LSB) in byte-order.

There exists an instruction, **thr**, to start executing new threads from desired code locations in memory.

## 2.1 Memory Organization

### Notes

System page size is 4096 bytes.

Address	Purpose	Brief
0	interrupt vector	interrupt handler descriptors
4096	keyboard buffer	keyboard input queue
8192	text segment	application program code (read-execute)
8192 + TEXTSIZE	data segment	program data (read-write)
DATA + DATASIZE	BSS segment	uninitialised data (runtime-allocated and zeroed)
MEMSIZE - 3.5 G	dynamic segment	free space for slab allocator
3.5 gigabytes	graphics	32-bit ARGB-format draw buffer

- the VM's 'physical' memory size is currently specified as **MEMSIZE**; this memory is mapped 1-to-1 to virtual address space
- thread stacks live at **MEMSIZE - thrid \* THRSTKSIZE**, i.e. at top of 'physical' address space

## Chapter 3

# Instruction Set

The ZPM instruction set was designed to resemble the C language closely, as well as to support a RISC-oriented set of typical machine operations.

### 3.1 Instruction Reference

#### 3.1.1 Instruction Set

##### Operands

The table below lists operand types.

- **i** stands for immediate operand
- **r** stands for register operand
- **m** stands for memory operand

##### Flags

Certain instructions set bits in the machine status word register (MSW). This is documented here on per-instruction basis.

- **z** stands for zero flag (ZF)
- **c** stands for carry flag (CF)
- **o** stands for overflow flag (OF)
- **s** stands for sign flag (SF)

**TODO:** stack/call conventions for certain instructions such as THR



## 3.1.2 Instruction Table

Mnemonic	Source	Destination	Brief	Flags
<b>not</b>	r	N/A	src = $\neg$ src;	z
<b>and</b>	r, i	r	dest = dest & src;	z
<b>or</b>	r, i	r	dest = dest   src;	N/A
<b>xor</b>	r, i	r	dest = dest $\oplus$ src;	z
<b>shl</b>	r, i	r	dest = dest $\ll$ src;	o, c
<b>shr</b>	r, i	r	dest = dest $\gg$ src; (zero-fill)	z
<b>sar</b>	r, i	r	dest = dest $\gg$ src; (sign-fill)	z
<b>rol</b>	r, i	r	dest = dest ROL src;	c
<b>ror</b>	r, i	r	dest = dest ROR src;	c
<b>inc</b>	r, i	N/A	src++;	o
<b>dec</b>	r, i	N/A	src--;	o, z
<b>add</b>	r, i	r	dest = dest + src;	o, z
<b>sub</b>	r, i	r	dest = dest - src;	s, z
<b>cmp</b>	r, i	r	compare two values and set flags	s, z
<b>mul</b>	r, i	r	dest = dest * src;	o, s, z
<b>div</b>	r, i	r	dest = dest / src;	s, z
<b>jmp</b>	r, i	branch to src	N/A	
<b>bz</b>	r, i	N/A	branch to src if (CF == 0)	N/A
<b>bnz</b>	r, i	N/A	branch to src if (CF != 0)	N/A
<b>blt</b>	r, i	N/A	branch to src if (SF != 0)	N/A
<b>ble</b>	r, i	N/A	branch to src if (SF != 0)    (ZF == 0)	N/A
<b>bgt</b>	r, i	N/A	branch to src if (SF != 0) && (ZF != 0)	N/A
<b>bge</b>	r, i	N/A	branch to src if (SF != 0)    (ZF == 0)	N/A
<b>bo</b>	r, i	N/A	branch to src if (OF != 0)	N/A
<b>bno</b>	r, i	N/A	branch to src if (OF == 0)	N/A
<b>bc</b>	r, i	N/A	branch to src if (CF != 0)	N/A
<b>bnc</b>	r, i	N/A	branch to src if (CF == 0)	N/A
<b>pop</b>	r	N/A	pop top of stack	N/A
<b>popa</b>	N/A	N/A	pop all registers from stack	N/A
<b>push</b>	r, i	N/A	push value on stack	N/A
<b>pusha</b>	N/A	N/A	push all register on stack	N/A
<b>lda</b>	r, i, m	r	load 32-bit longword	N/A
<b>sta</b>	r	r, i, m	store 32-bit longword	N/A
<b>call</b>	r, i	N/A	call subroutine; construct stack frame	N/A
<b>enter</b>	N/A	N/A	enter subroutine	N/A
<b>leave</b>	N/A	N/A	leave subroutine	N/A
<b>ret</b>	N/A	N/A	return from subroutine	N/A
<b>thr</b>	m	N/A	start new thread at address	N/A
<b>ltb</b>	r, i	N/A	load thread-local storage base address	N/A
<b>ldr</b>	r, m	r	load special register	N/A
<b>str</b>	r	r, m	store special register	N/A
<b>rst</b>	m	N/A	reset machine	N/A
<b>hlt</b>	m	N/A	halt machine	N/A
<b>inb</b>	r, i	N/A	read 8-bit byte from input port	N/A
<b>inw</b>	r, i	N/A	read 16-bit word from input port	N/A
<b>inl</b>	r, i	N/A	read 32-bit longword from input port	N/A
<b>outb</b>	r, i	N/A	write 8-bit byte to input port	N/A
<b>outw</b>	r, i	N/A	write 16-bit word to input port	N/A
<b>outl</b>	r, i	N/A	write 32-bit longword to input port	N/A

### 3.1.3 I/O Port Map

#### I/O Address Space

The I/O address space maps 65,536 I/O ports to unsigned 16-bit address space.

Port #	Name	Default
<b>0x0000</b>	STDIN	keyboard input
<b>0x0001</b>	STDOUT	console output
<b>0x0002</b>	STDERR	error output



# Chapter 4

## Assembly

### 4.1 Syntax

#### AT&T Syntax

We use so-called AT&T-syntax assembly. Perhaps the most notorious difference from Intel-syntax is the operand order; AT&T lists the source operand first, destination second, whereas Intel syntax does it vice versa.

#### Symbol Names

Label names must start with an underscore or a letter; after that, the name may contain underscores, letters, and digits. Label names end with a ':', so like

```
value:      .long 0xb4b5b6b7
```

would declare a longword value at the address of **value**.

#### Instructions

The instruction operand order is source first, then destination. For example,

```
lda        8(%r0), %r1
```

would load the value from address **r0 + 8** to the register **r1**.

#### Operands

Register operand names are prefixed with a '%'. Immediate constants and direct addresses are prefixed with a '#'. Label addresses are referred to as their names without prefixes.

The assembler supports simple preprocessing (of constant-value expressions), so it is possible to do things such as

```
.define     FLAG1      0x01
.define     FLAG2      0x02

lda        $(FLAG1|FLAG2), %r1
```

### Registers

Register names are prefixed with '%'; there are 16 registers r0..r15. For example,

```
add    %r0, %r1
```

would add the longword in r0 to r1.

### Direct Addressing

Direct addressing takes the syntax

```
lda    val, %r0
```

which moves the longword at **address val** into r0.

### Indexed Addressing

Indexed addressing takes the syntax

```
lda    4(%r0), %r1
```

where 4 is an integral constant offset and r0 is a register name. In short, this would store the value at the address **r0 + 4** into r1.

### Indirect Addressing

Indirect addresses are indicated with a '\*', so

```
lda    *%r0, %r1
```

would store the value from the **address in the register r0** into register r1, whereas

```
lda    *val, %r0
```

would move the value **pointed to by val** into r0.

Note that the first example above was functionally equivalent with

```
lda    (%r0), %r1
```

### Immediate Addressing

Immediate addressing takes the syntax

```
lda    $str, %r0
```

which would store the **address of str** into r0.

## 4.2 Assembler Directives

## 4.3 Input Directives

### 4.3.1 .include

The **.include** directive takes the syntax

```
.include <stdio.inc>
```

to insert `<stdio.inc>` into the translation stream verbatim.

### 4.3.2 `.import`

The `.import` directive takes the syntax

```
.import <file.asm>
```

or

```
.import <file.obj>
```

to import foreign assembly or object files into the stream. **Note** that only symbols declared with `.globl` will be made globally visible to avoid namespace pollution.

## 4.4 Link Directives

### 4.4.1 `.org`

The `.org` directive takes a single argument and sets the linker location address to the given value.

### 4.4.2 `.space`

The `.space` directive takes a single argument and advances the link location address by the given value.

### 4.4.3 `.align`

The `.align` directive takes a single argument and aligns the next label, data, or instruction to a boundary of the given size.

### 4.4.4 `.globl`

The `.globl` directive takes one or several symbol names arguments and declares the symbols to have global visibility (linkage).

## 4.5 Data Directives

### 4.5.1 `.long`

`.long` takes any number of arguments and declares in-memory 32-bit entities.

### 4.5.2 .byte

**.byte** takes any number of arguments and declares in-memory 8-bit entities.

### 4.5.3 .short

**.short** takes any number of arguments and declares in-memory 16-bit entities.

### 4.5.4 .asciz

**.asciz** takes a C-style string argument of characters enclosed within double quotes ("). Escape sequences '\n' (newline), '\t' (tabulator), and '\r' (carriage return) are supported.

### 4.5.5 Preprocessor Directives

## 4.6 .define

**.define** lets one declare symbolic names for constant (numeric) values. For example, if you have

```
<hook.def>
```

```
.define STDIN 0
.define STDOUT 1
.define STDERR 2
```

## 4.7 Input and Output

The pseudo machine uses some predefined ports for keyboard and console I/O. The currently predefined ports are

Port	Use	Notes
0x00	keyboard input	interrupt-driven
0x01	console output	byte stream
0x02	error output	directed to console by default

## 4.8 Simple Program

The following code snippet prints the string "hello" a newline to the console. Note that the string is saved using the standard C convention of NUL-character termination.

```

msg:      .asciz      "hello\n"

.align    4

_start:
    sta     $msg, %r0
    ldb     *%r0, %r1
    ldb     $0x01, %r2
    cmp     $0x00, %r1
    bz      done
loop:
    inc     %r0
    outb    %r1, %r2
    ldb     *%r0, %r1
    cmp     $0x00, %r1
    bnz     loop
done:
    hlt

```

## 4.9 Threads

The pseudo machine supports hardware threads with the **thr** instruction. It takes a single argument, which specifies the new execution start address; function arguments should be passed in registers.

### 4.10 Example Program

The following piece of code shows simple utilisation of threads.

```

.import <bzero.asm>

memzero:
    lda     $65536, %r0    // address
    lda     $4096, %r1     // length
    call    bzero
    hlt

_start:
    thr     $memzero
    hlt

```

### 4.11 Interrupts

Software- and CPU-generated interrupts are often referred to as **traps**. I call those and hardware-generated **interrupt requests** interrupts, collectively.

### 4.11.1 Interrupt Interface

The lowest page (4096 bytes) in virtual machine address space contains the **interrupt vector**, i.e. a table of interrupt handler addresses to trigger them.

Interrupt handler invocations only push the **program counter** and **old frame pointer**, so you need to reserve the registers you use manually. This is so interrupts could be as little overhead as possible to handle.

### 4.11.2 Keyboard Input

In order to read keyboard input without polling, we need to hook the **interrupt 0**. This is done in two code modules; an interrupt handler as well as other support code.

I will illustrate the interrupt handler first.

### 4.11.3 Keyboard Interrupt Handler

**TODO:** example interrupt handler

#### 4.11.3.1 Keyboard Support Code

**TODO:** queue keypresses in 16-bit values; 32-bit if full Unicode requested.